

# NoSQL & Key-Value Databases Notes

## Distributed Databases and Concurrency Models

### Pessimistic Concurrency (ACID)

- Focuses on "data safety" and assumes conflicts will occur
- Transactions protect themselves from other transactions
- Uses locking resources (both read and write locks) until a transaction is complete
- Similar to borrowing a book from a library - if you have it, no one else can use it
- Prioritizes data consistency and integrity

### Optimistic Concurrency

- Does not obtain locks when reading or writing data
- Assumes conflicts are unlikely to occur
- Implements validation through timestamps and version numbers
- When a transaction is ready to commit, it checks if data has been modified by other transactions
- If a conflict is detected, the transaction can be rolled back and re-executed
- Works well in low-conflict systems (read-heavy workloads, analytical databases, backups)
- Less efficient in high-conflict systems due to the overhead of rolling back transactions

## CAP Theorem

- States that distributed systems can have at most two of the following three properties:
  1. **Consistency**: Every user has an identical view of data at any moment
  2. **Availability**: System remains operational during failures
  3. **Partition Tolerance**: System can maintain operations when network partitions occur
- The three possible combinations:
  1. **CA (Consistency + Availability)**: Systems always respond with the latest data, but may not handle network partitions well
  2. **CP (Consistency + Partition Tolerance)**: If the system responds, it's with the latest data, but requests may be dropped during partitions
  3. **AP (Availability + Partition Tolerance)**: System always responds, but may not return the absolute latest data

- Traditional relational databases (PostgreSQL, MySQL) typically prioritize CA
- Many NoSQL systems (MongoDB, Redis) prioritize CP
- Some distributed systems (Cassandra, DynamoDB) prioritize AP

## BASE - Alternative to ACID for Distributed Systems

- **Basically Available:** System guarantees availability but responses might be inconsistent or in a changing state
- **Soft State:** System state may change over time without input due to eventual consistency
- **Eventual Consistency:** System will eventually become consistent once all updates stop and are propagated

## NoSQL Databases

- Term "NoSQL" originated in 1998 but has evolved to mean "Not Only SQL"
- Often refers to non-relational database systems
- Developed partly as a response to handling unstructured web-based data
- Allows for more flexible data models than traditional relational databases

## Key-Value Databases

### Characteristics

- **Simplicity:** Extremely simple data model compared to relational tables
- **Speed:** Usually deployed as in-memory databases
- **O(1) Operations:** Retrieving values by key is typically constant time
- **No Complex Queries:** No support for joins or complex queries
- **Scalability:** Horizontal scaling is simple by adding more nodes
- **Eventual Consistency:** In distributed environments, nodes will eventually converge on the same value

### Use Cases

- **Data Science Applications:**
  - Storing EDA/experimentation results
  - Feature stores for low-latency ML model training and inference
  - Model monitoring metrics
- **Software Engineering Applications:**

- Session information storage
- User profiles and preferences
- Shopping cart data
- Caching layer in front of slower disk-based databases

## Redis (Remote Dictionary Server)

- Open-source, in-memory data structure store
- Primarily a key-value store but supports multiple data models
- Consistently ranked as the #1 key-value store according to db-engines.com
- Developed in 2009 in C++
- Can perform over 100,000 SET operations per second
- Supports durability through snapshots and append-only files
- Organized into 16 logical databases (numbered 0-15)

## Redis Data Types

### 1. Strings

- Simplest data type, maps a string key to a string value
- Can store text, integers, serialized objects, binary arrays
- Use cases: Caching, config settings, counters, rate limiting
- Commands: SET, GET, INCR, DECR, INCRBY, DECRBY, SETNX

### 2. Hashes

- Collection of field-value pairs under a single key
- Maximum of  $2^{32}-1$  field-value pairs per hash
- Use cases: Representing objects, session management, user tracking
- Commands: HSET, HGET, HGETALL, HMGET, HINCRBY, HKEYS

### 3. Lists

- Linked lists of string values
- $O(1)$  operations for insertion at front or end
- Use cases: Stacks, queues, message passing, social media feeds, logging
- Commands:
  - Queue operations: LPUSH, RPOP
  - Stack operations: LPUSH, LPOP
  - Other operations: LLEN, LRANGE

### 4. Sets

- Unordered collections of unique strings
- Support for powerful set operations (union, intersection, difference)
- Use cases: Tracking unique items, access control, group membership

- Commands: SADD, SISEMEMBER, SCARD, SINTER, SDIFF, SREM, SRANDMEMBER

## 5. JSON

- Full support for the JSON standard
- Uses JSONPath syntax for navigation
- Stored in binary as a tree structure for fast access to elements

## Redis Setup

- Can be run locally via Docker
- Default port is 6379
- Requires security considerations in production:
  - Should not expose Redis port publicly
  - Should set authentication passwords
  - Should consider encryption for sensitive data